

STA 445: Assignment 4

Kaylin McLiverty

2023-11-03

Chapter 13

Exercise 1

A common task is to take a set of data that has multiple categorical variables and create a table of the number of cases for each combination. An introductory statistics textbook contains a dataset summarizing student surveys from several sections of an intro class. The two variables of interest for us are **Gender** and **Year** which are the students gender and year in college.

- a) Download the dataset and correctly order the `Year` variable using the following: `Survey <- read.csv('https://www.lock5stat.com/datasets3e/StudentSurvey.csv', na.strings=c('', ' '))`

```
# Exercise 1 part a
# Downloading the dataset
Survey <- read.csv('https://www.lock5stat.com/datasets3e/StudentSurvey.csv',
  na.strings=c('', ' '))
head(Survey)
```

	Year	Sex	Smoke	Award	HigherSAT	Exercise	TV	Height	Weight	Siblings
## 1	Senior	M	No	Olympic	Math	10	1	71	180	4
## 2	Sophomore	F	Yes	Academy	Math	4	7	66	120	2
## 3	FirstYear	M	No	Nobel	Math	14	5	72	208	2
## 4	Junior	M	No	Nobel	Math	3	1	63	110	1
## 5	Sophomore	F	No	Nobel	Verbal	3	3	65	150	1
## 6	Sophomore	F	No	Nobel	Verbal	5	4	65	114	2

	BirthOrder	VerbalSAT	MathSAT	SAT	GPA	Pulse	Piercings
## 1	4	540	670	1210	3.13	54	0
## 2	2	520	630	1150	2.50	66	3
## 3	1	550	560	1110	2.55	130	0
## 4	1	490	630	1120	3.10	78	0
## 5	1	720	450	1170	2.70	40	6
## 6	2	600	550	1150	3.20	80	4

```
#don't have to worry about ordering
#order(Survey$Year, decreasing = FALSE, method=c('FirstYear', 'Sophomore', 'Junior', 'Senior'))
```

- b) Using some combination of `dplyr` functions, produce a data set with eight rows that contains the number of responses for each gender:year combination. Make sure your table orders the `Year` variable in the correct order of `First Year`, `Sophomore`, `Junior`, and then `Senior`. *You might want to look at the following functions: `dplyr::count` and *`dplyr::drop_na`.*

```
# Exercise 1 part b
survey1 <- Survey %>%
  count(Sex,Year) %>%
  drop_na()
survey1
```

```
##   Sex      Year  n
## 1  F FirstYear 43
## 2  F   Junior 18
## 3  F   Senior 10
## 4  F Sophomore 96
## 5  M FirstYear 51
## 6  M   Junior 17
## 7  M   Senior 26
## 8  M Sophomore 99
```

c) Using `tidyr` commands, produce a table of the number of responses in the following form:

Gender	First Year	Sophomore	Junior	Senior
Female				
Male				

```
# Exercise 1 part c
survey2 <- survey1 %>% pivot_wider(
  names_from=Year,
  values_from=n)
survey2
```

```
## # A tibble: 2 x 5
##   Sex   FirstYear Junior Senior Sophomore
##   <chr>      <int>  <int>  <int>      <int>
## 1 F           43     18    10         96
## 2 M           51     17    26         99
```

Exercise 2

From the book website, there is a .csv file of the daily maximum temperature in Flagstaff at the Pulliam Airport. The direction link is at: <https://raw.githubusercontent.com/dereksonderegger/444/master/data-raw/FlagMaxTemp.csv>

a) Create a line graph that gives the daily maximum temperature for 2005.
Make sure the x-axis is a date and covers the whole year.

```
# Exercise 2 part a
#loading the .csv file
maxtemp <- read.csv('https://raw.githubusercontent.com/dereksonderegger/444/master/data-raw/FlagMaxTemp.csv')
maxtemp_05 <- filter(maxtemp, Year==2005)

maxtemp_05 <- pivot_longer(
  maxtemp_05,
  X1:X31,
  names_to='Day',
  values_to='Temperature')
```

```
maxtemp_05 <- unite(maxtemp_05, col='Date', c('Year','Month','Day'))
```

```
maxtemp_05 <- maxtemp_05 %>%
  mutate(date=lubridate::ymd(Date)) %>% drop_na()
```

```
## Warning: There was 1 warning in `mutate()`.
## i In argument: `date = lubridate::ymd(Date)`.
```

```
## Caused by warning:
```

```
## ! 7 failed to parse.
```


```
head(maxtemp_05)
```

```
## # A tibble: 6 x 4
```

	X	Date	Temperature	date
	<int>	<chr>	<dbl>	<date>
## 1	235	2005_1_X1	37.9	2005-01-01
## 2	235	2005_1_X2	37.9	2005-01-02
## 3	235	2005_1_X3	39.0	2005-01-03
## 4	235	2005_1_X4	34.0	2005-01-04
## 5	235	2005_1_X5	35.1	2005-01-05
## 6	235	2005_1_X6	27.0	2005-01-06

```
#Creating a line graph
```

```
ggplot(maxtemp_05, aes(x=date, y=Temperature )) +
  geom_line()
```



- b) Create a line graph that gives the monthly average maximum temperature for 2013 - 2015. *Again the x-axis should be the date and the axis*

```

    *spans 3 years.*
# Exercise 2 part b
maxtemp_35 <- filter(maxtemp, Year >= 2013)

maxtemp_35 <- maxtemp_35 %>% group_by(Month, Year)

maxtemp_35 <- pivot_longer(
  maxtemp_35,
  X1:X31,
  names_to='day',
  values_to='Temperature') %>% drop_na()
maxtemp_35

## # A tibble: 1,001 x 5
## # Groups:   Month, Year [35]
##       X   Year Month day   Temperature
##   <int> <int> <int> <chr>     <dbl>
## 1    330  2013     1 X1         30.0
## 2    330  2013     1 X2         28.9
## 3    330  2013     1 X3         36.0
## 4    330  2013     1 X4         48.0
## 5    330  2013     1 X5         43.0
## 6    330  2013     1 X6          41
## 7    330  2013     1 X7         46.0
## 8    330  2013     1 X8         46.9
## 9    330  2013     1 X9         46.9
## 10   330  2013     1 X10        46.9
## # i 991 more rows

maxtemp_35 <- maxtemp_35 %>%
  group_by(Month, Year) %>%
  summarise(MeanTemp=mean(Temperature)) %>%
  mutate(Day=1)

## `summarise()` has grouped output by 'Month'. You can override using the
## `.groups` argument.

maxtemp_35

## # A tibble: 35 x 4
## # Groups:   Month [12]
##   Month   Year MeanTemp   Day
##   <int> <int>     <dbl> <dbl>
## 1     1  2013     41.2     1
## 2     1  2014     48.6     1
## 3     1  2015     48.8     1
## 4     2  2013     42.7     1
## 5     2  2014     47.9     1
## 6     2  2015     53.9     1
## 7     3  2013     56.9     1
## 8     3  2015     56.7     1
## 9     4  2013     61.6     1
## 10    4  2014     57.9     1
## # i 25 more rows


```

```
maxtemp_35 <- unite(maxtemp_35, col='Date', c('Year','Month','Day')) %>%
  mutate(Date=lubridate::ymd(Date))
```

```
maxtemp_35
```

```
## # A tibble: 35 x 2
##   Date      MeanTemp
##   <date>     <dbl>
## 1 2013-01-01    41.2
## 2 2014-01-01    48.6
## 3 2015-01-01    48.8
## 4 2013-02-01    42.7
## 5 2014-02-01    47.9
## 6 2015-02-01    53.9
## 7 2013-03-01    56.9
## 8 2015-03-01    56.7
## 9 2013-04-01    61.6
## 10 2014-04-01    57.9
## # i 25 more rows
```

```
#Creating a line graph
ggplot(maxtemp_35, aes(x=Date, y=MeanTemp )) +
  geom_line()
```



Exercise 4

For this problem we will consider two simple data sets.

#Code given for Exercise 4:

```
A <- tribble(
  ~Name, ~Car,
  'Alice', 'Ford F150',
  'Bob', 'Tesla Model III',
  'Charlie', 'VW Bug')

B <- tribble(
  ~First.Name, ~Pet,
  'Bob', 'Cat',
  'Charlie', 'Dog',
  'Alice', 'Rabbit')
```

a) Squish the data frames together to generate a data set with three rows and three columns. Do two ways: first using ``cbind`` and then using one of the ``dplyr`` ``join`` commands.

Exercise 4 part a

#Using cbind

```
B2 <- B[order(B$First.Name), ]
B2
```

```
## # A tibble: 3 x 2
##   First.Name Pet
##   <chr>      <chr>
## 1 Alice      Rabbit
## 2 Bob        Cat
## 3 Charlie   Dog
```

```
cbind(A, Pet=B2$Pet)
```

```
##      Name      Car      Pet
## 1  Alice    Ford F150 Rabbit
## 2    Bob Tesla Model III  Cat
## 3 Charlie   VW Bug    Dog
```

#using 'dplyr' 'join' commands

```
inner_join(A,B2, by=c('Name'='First.Name'))
```

```
## # A tibble: 3 x 3
##   Name      Car      Pet
##   <chr>  <chr>    <chr>
## 1 Alice  Ford F150  Rabbit
## 2 Bob    Tesla Model III Cat
## 3 Charlie VW Bug    Dog
```

b) It turns out that Alice also has a pet guinea pig. Add another row to the ``B`` data set. Do this using either the base function ``rbind``, or either of the ``dplyr`` functions ``add_row`` or ``bind_rows``.

Exercise 4 part b

#Used cbind

```
pig <- c('Alice', 'Guinea Pig')
```

```
B3 <- rbind(B2,pig)
B3
```

```
## # A tibble: 4 x 2
##   First.Name Pet
##   <chr>      <chr>
## 1 Alice      Rabbit
## 2 Bob        Cat
## 3 Charlie    Dog
## 4 Alice      Guinea Pig
```

c) Squish the `A` and `B` data sets together to generate a data set with four rows and three columns. Do this two ways: first using `cbind` and then using one of the `dplyr` `join` commands. Which was easier to program? Which is more likely to have an error.

```
# Exercise 4 part c
#Using cbind
Q <- cbind(A, Pet=B2$Pet)
Q <- rbind(Q,pig)
```

```
## Warning in rbind(deparse.level, ...): number of columns of result, 3, is not a
## multiple of vector length 2 of arg 2
```

```
Q
```

```
##      Name      Car      Pet
## 1  Alice    Ford F150 Rabbit
## 2    Bob Tesla Model III  Cat
## 3 Charlie      VW Bug   Dog
## 4  Alice    Guinea Pig  Alice
```

```
#using 'dplyr' 'join' commands
R <- full_join(A,B3, by=c('Name'='First.Name'))
R
```

```
## # A tibble: 4 x 3
##   Name      Car      Pet
##   <chr>   <chr>   <chr>
## 1 Alice  Ford F150  Rabbit
## 2 Alice  Ford F150  Guinea Pig
## 3 Bob    Tesla Model III Cat
## 4 Charlie VW Bug    Dog
```

The second way of using one of the `dplyr` `join` commands was easier to program and is less likely to have error. Using `cbind` is bound for more error and the function itself is a bit more finicky to work with. I found it easier to combine the columns first and then add the additional row

Note for Exercise 5: Respectfully, I have looked over Exercise 5 and will review it in more detail. I know that this exercise is difficult and will not be graded. Thus, I have not made a productive attempt yet.

Exercise 5

Data table joins are extremely common because effective database design almost always involves having multiple tables for different types of objects. To illustrate both the table joins and the usefulness of multiple tables we will develop a set of data frames that will represent a credit card company's customer data base. We will have tables for Customers, Retailers, Cards, and Transactions.

```
# Given code for Exercise 5
Customers <- tribble(
  ~PersonID, ~Name, ~Street, ~City, ~State,
```

```

1, 'Derek Sonderegger', '231 River Run', 'Flagstaff', 'AZ',
2, 'Aubrey Sonderegger', '231 River Run', 'Flagstaff', 'AZ',
3, 'Robert Buscaglia', '754 Forest Heights', 'Flagstaff', 'AZ',
4, 'Roy St Laurent', '845 Elk View', 'Flagstaff', 'AZ')

Retailers <- tribble(
  ~RetailID, ~Name, ~Street, ~City, ~State,
  1, 'Kickstand Kafe', '719 N Humphreys St', 'Flagstaff', 'AZ',
  2, 'MartAnnes', '112 E Route 66', 'Flagstaff', 'AZ',
  3, 'REI', '323 S Windsor Ln', 'Flagstaff', 'AZ' )

Cards <- tribble(
  ~CardID, ~PersonID, ~Issue_DateTime, ~Exp_DateTime,
  '9876768717278723', 1, '2019-9-20 0:00:00', '2022-9-20 0:00:00',
  '5628927579821287', 2, '2019-9-20 0:00:00', '2022-9-20 0:00:00',
  '7295825498122734', 3, '2019-9-28 0:00:00', '2022-9-28 0:00:00',
  '8723768965231926', 4, '2019-9-30 0:00:00', '2022-9-30 0:00:00' )

Transactions <- tribble(
  ~CardID, ~RetailID, ~DateTime, ~Amount,
  '9876768717278723', 1, '2019-10-1 8:31:23', 5.68,
  '7295825498122734', 2, '2019-10-1 12:45:45', 25.67,
  '9876768717278723', 1, '2019-10-2 8:26:31', 5.68,
  '9876768717278723', 1, '2019-10-2 8:30:09', 9.23,
  '5628927579821287', 3, '2019-10-5 18:58:57', 68.54,
  '7295825498122734', 2, '2019-10-5 12:39:26', 31.84,
  '8723768965231926', 2, '2019-10-10 19:02:20', 42.83)

Cards <- Cards %>%
  mutate( Issue_DateTime = lubridate::ymd_hms(Issue_DateTime),
           Exp_DateTime = lubridate::ymd_hms(Exp_DateTime) )
Transactions <- Transactions %>%
  mutate( DateTime = lubridate::ymd_hms(DateTime))

```

- a) Create a table that gives the credit card statement for Derek. It should give all the transactions, the amounts, and the store name. Write your code as if the only initial information you have is the customer's name.
 Hint: Do a bunch of table joins, and then filter for the desired customer
 name. To be efficient, do the filtering first and then do the table joins.

Exercise 5 part a

- b) Aubrey has lost her credit card on Oct 15, 2019. Close her credit card at 4:28:21 PM and issue her a new credit card in the `Cards` table.
 Hint: Using the Aubrey's name, get necessary CardID and PersonID and save
 those as `cardID` and `personID`. Then update the `Cards` table row that
 corresponds to the `cardID` so that the expiration date is set to the time
 that the card is closed. Then insert a new row with the `personID` for
 Aubrey and a new `CardID` number that you make up.

Exercise 5 part b

- c) Aubrey is using her new card at Kickstand Kafe on Oct 16, 2019 at 2:30:21 PM for coffee with a charge of \$4.98. Generate a new transaction

for this action.
Hint: create temporary variables `card`,`retailid`,`datetime`, and
`amount` that contain the information for this transaction and then
write your code to use those. This way in the next question you can just
use the same code but modify the temporary variables. Alternatively, you
could write a function that takes in these four values and manipulates the
tables in the GLOBAL environment using the `<<-` command to assign a result
to a variable defined in the global environment. The reason this is OK is
that in a real situation, these data would be stored in a database and we
would expect the function to update that database.

Exercise 5 part c

- d) On Oct 17, 2019, some nefarious person is trying to use her OLD credit card at REI. Make sure your code in part (c) first checks to see if the credit card is active before creating a new transaction. Using the same code, verify that the nefarious transaction at REI is denied.

Exercise 5 part d

- e) Generate a table that gives the credit card statement for Aubrey. It should give all the transactions, amounts, and retailer name for both credit cards she had during this period.

Exercise 5 part e